

# UPC

**GIT repository :** [https://github.com/AurelienVasseur/LO47\\_UPC\\_HandsOn](https://github.com/AurelienVasseur/LO47_UPC_HandsOn)

---

## 3 - Simplified 1D Laplace solver

### 3.1 - A simplified 1-D Laplace solver in C

This first 1-D Laplace solver in C is not optimized. Indeed, We don't use the threads power to increase the program performance. The program do all iterations. In fact, there is no work sharing between different threads.

### 3.2 - The 1D solver in UPC

In this solution we are going to share work between many threads. To do this, we start by create shared variables in the shared space. These variables will be used by all of the threads.

```
8  //== declare the x, x_new, b arrays in the shared space with size of TOTALSIZE
9  shared double x[TOTALSIZE];
10 shared double x_new[TOTALSIZE];
11 shared double b[TOTALSIZE];
12
13 void init();
14
15 int main(int argc, char **argv){
16     int j;
17
18     init();
19     upc_barrier;
20
21     //== add a for loop which goes through the elements in the x_new array
22     for( j=1; j<(TOTALSIZE)-1; j++ ){
23         //== insert an if statement to do the work sharing across the threads
24         if( (j%THREADS) == MYTHREAD){
25             x_new[j] = 0.5*( x[j-1] + x[j+1] + b[j] );
26         }
27     }
```

To share the work we divide it inside the loop which goes through the elements in the `x_new` array. To do this, we add a if statement which will allow each thread to test if it is up to him to do this task. This is a first way to do the work sharing.

```

21 //== add a for loop which goes through the elements in the x_new array
22 for( j=1; j<(TOTALSIZE)-1; j++ ){
23     //== insert an if statement to do the work sharing across the threads
24     if( (j%THREADS) == MYTHREAD){
25         x_new[j] = 0.5*( x[j-1] + x[j+1] + b[j] );
26     }
27 }
28
29 if( MYTHREAD == 0 ){
30     printf("    b    |    x    | x_new\n");
31     printf("===== \n");
32
33     for( j=0; j<TOTALSIZE; j++ )
34         printf("%1.4f | %1.4f | %1.4f \n", b[j], x[j], x_new[j]);
35 }

```

After that, we start printing the result on the screen. To do this work we will only use the thread 0. This is why there is a if statement which check if it is thread 0.

With this solution we will meet a problem. Thread 0 may start printing result before others finish. The problem is that there is no condition which force threads to expect them between the end of the main work (l. 27) and the start of the printing (l. 29). If we want to start printing result only after all threads finished their work we need to add a `upc_barrier`.

```

21 //== add a for loop which goes through the elements in the x_new array
22 for( j=1; j<(TOTALSIZE)-1; j++ ){
23     //== insert an if statement to do the work sharing across the threads
24     if( (j%THREADS) == MYTHREAD){
25         x_new[j] = 0.5*( x[j-1] + x[j+1] + b[j] );
26     }
27 }
28
29 upc_barrier;
30
31 if( MYTHREAD == 0 ){
32     printf("    b    |    x    | x_new\n");
33     printf("===== \n");
34
35     for( j=0; j<TOTALSIZE; j++ )
36         printf("%1.4f | %1.4f | %1.4f \n", b[j], x[j], x_new[j]);
37 }

```

### 3.3 - Better work sharing construct with a single for loop

In this solution we will improve the for loop to perform the work distribution across threads.

```

20 // ==> setup j to point to the first element so that the current thread should
21 // progress (in respect to its affinity)
22 j = MYTHREAD;
23
24 // ==> add a for loop which goes only through the elements in the x_new array
25 // with affinity to the current THREAD
26
27 for( j=MYTHREAD+1; j<(TOTALSIZE*THREADS)-1; j+=THREADS )
28     x_new[j] = 0.5*( x[j-1] + x[j+1] + b[j] );
29
30 upc_barrier;
31
32 if( MYTHREAD == 0 ){
33     printf("    b    |    x    | x_new\n");
34     printf("===== \n");
35
36     for( j=0; j<TOTALSIZE*THREADS; j++ )
37         printf("%1.4f | %1.4f | %1.4f \n", b[j], x[j], x_new[j]);
38 }

```

To do that we will edit the for loop which was coupled with a if statement in the previous solution. In fact we will delete the if statement and update the loop. In the previous solution each thread do each iteration of the loop (which correspond of the array size) even if it was to him to do the work. Now, each thread do only iterations which match with his work.

### 3.4 - Blocked arrays and work sharing with upc\_forall

In this new solution we will use the upc\_forall power which help us to share work between many threads. In fact, upc\_forall allows the distribution of the iterations according to the affinity.

```

6  #define BLOCKSIZE 16
7
8  //==> declare the x, x_new and b arrays in the shared space with size of
9  //    BLOCKSIZE*THREADS and with blocking size of BLOCKSIZE
10 shared [BLOCKSIZE] double x[BLOCKSIZE*THREADS];
11 shared [BLOCKSIZE] double x_new[BLOCKSIZE*THREADS];
12 shared [BLOCKSIZE] double b[BLOCKSIZE*THREADS];
13
14 void init();
15
16 int main(int argc, char **argv){
17     int j;
18
19     init();
20     upc_barrier;
21
22     //==> insert a upc_forall statement to do work sharing while
23     //    respecting the affinity of the x_new array
24     upc_forall( j=1; j<(BLOCKSIZE*THREADS)-1; j++; j){
25         x_new[j] = 0.5*( x[j-1] + x[j+1] + b[j] );
26     }
27     upc_barrier;

```

### 3.5 - Synchronization

```

20 // add two barrier statements, to ensure all threads finished computing
21 // x_new[] and to ensure that all threads have completed the array
22 // swapping.
23 for( iter=0; iter<10000; iter++ ){
24     upc_forall( j=1; j<BLOCKSIZE*THREADS-1; j++; &x_new[j] ){
25         x_new[j] = 0.5*( x[j-1] + x[j+1] + b[j] );
26     }
27     upc_barrier;
28     upc_forall( j=0; j<BLOCKSIZE*THREADS; j++; &x_new[j] ){
29         x[j] = x_new[j];
30     }
31     upc_barrier;
32 }

```

In this new example we want to synchronize all of the threads. To do that we added two barrier statements which force threads to expect them. These barrier are two `upc_barrier` (l.27 & l.31).

### 3.6 - Reduction operation

In this last example we want to keep track of the maximum difference between the  $x$  and  $x_{\text{new}}$ . To begin, each thread need to compute a local version of the difference. After that, we need to compare all of these values to find which is the biggest.

```
41     upc_forall( j=1; j<TOTALSIZE*THREADS-1; j++; &x_new[j] ){
42         x_new[j] = 0.5 * ( x[j-1] + x[j+1] + b[j] );
43
44         if( diff[MYTHREAD] < x_new[j] - x[j] )
45             diff[MYTHREAD] = x_new[j] - x[j];
46     }
47
48     // Each thread as a local value for diff
49     // The maximum of those values should be used to check
50     // the convergence.
51     diffmax = diff[0];
52     for(i=1; i<THREADS; i++)
53     {
54         if(diffmax < diff[i])
55             diffmax = diff[i];
56     }
```

Indeed, we need to browse each thread diff value. To do that we use a for loop (l.52). Inside the loop we check with a if statement if the diff value of the thread is the biggest. If is true we save this value on diffmax.



## 4 - 2D Heat conduction

### 4.4 - First UPC program

In this part I created the UPC version of the 2D Heat conduction program. The following program is not optimized, this is only a conversion from C language to UPC.

```
1  #include <stdio.h>
2  #include <math.h>
3  #include <sys/time.h>
4  #include <upc_relaxed.h>
5
6  #define N 30
7  #define BLOCKSIZE N
8
9  shared[BLOCKSIZE] double grids[2][N][N][N];
10 shared double dTmax_local[THREADS];
11
12 void initialize(void)
13 {
14     int y, x;
15     for(y=1; y<N-1; y++) {
16         upc_forall(x=1; x<N-1; x++; &grid[0][0][y][x]) {
17             grids[0][0][y][x] = grids[1][0][y][x] = 1.0;
18         }
19     }
20 }
21
22 int main(void)
23 {
24     double dTmax, dT, epsilon;
25     int finished, z, y, x, i;
26     double T;
27     int nr_iter;
28     int sg, dg;
29
30     initialize();
```

```
32  /* set the constants */
33  epsilon = 0.0001;
34  finished = 0;
35  nr_iter = 0;
36  sg = 0;
37  dg = 1;
38
39  /* synchronization */
40  upc_barrier;
41
42  do
43  {
44      dTmax = 0.0;
45      for(z=1; z<N-1; z++) {
46          for(y=1; y<N-1; y++) {
47              upc_forall(x=1; x<N-1; x++: &grids[sg][z][y][x]) {
48                  T = (grids[sg][z+1][y][x] + grids[sg][z-1][y][x]
49                      + grids[sg][z][y-1][x] + grids[sg][z][y+1][x]
50                      + grids[sg][z][y][x-1] + grids[sg][z][y][x+1]) / 6.0;
51                  dT = T - grids[sg][z][y][x];
52                  grids[dg][z][y][x] = T;
53                  if (dTmax < fabs(dT))
54                  {
55                      dTmax = fabs(dT);
56                  }
57              }
58          }
59      }
60
61      dTmax_local[MYTHREAD] = dTmax;
62      upc_barrier;
```

```

63
64     dTmax = dTmax_local[0];
65     for(i=1; i<THREADS; i++) {
66         if (dTmax < dTmax_local[i]){
67             dTmax = dTmax_local[i];
68         }
69     }
70
71     upc_barrier;
72
73     if(dTmax < epsilon) {
74         finished = 1;
75     } else {
76         dg = sg;
77         sg = !sg;
78     }
79     nt_iter++;
80     while (!finished);
81
82     upc_barrier;
83
84     if(MYTHREAD == 0) {
85         printf("%d iterations \n", nr_iter);
86     }
87
88     return 0;
89 }
90

```

#### 4.5 - Better memory use

In this second version of the program, I optimized the memory use. Indeed, the previous version is copying the complete `grid_new[]` to `grid[]` at the end of each iteration. This step creates an overhead and can be avoided by implementing a pointer flipping. This is what I implemented in this new program.

Now, the grids `grid[]` and `grid_new[]` will be accessed using either `ptr[]` or `new_ptr[]`, each respectively pointing to `grid[]` and `grid_new[]`. At the end of an iteration, instead of copying the complete `grid_new[]` to `grid[]` (like the previous version), the pointers `ptr[]` and `new_ptr[]` are swapped (l. 111-116) and a new iteration can take place.



```
1  #include <stdio.h>
2  #include <math.h>
3  #include <sys/time.h>
4  #include <upc relaxed.h>
5
6  #define N 30
7
8  /* ----- */
9  /* Complete the grids declarations to be shared and blocked as the
10 | biggest chunk of rows */
11 shared [N+2] double grid[N+2][N+2], new_grid[N+2][N+2];
12 /* add a new shared array, in which each thread will maintain its local
13 | execution time */
14 shared double exectime[THREADS];
15 /* now declare dTmax_local, array of THREADS elements, as shared. This
16 | array is used to reduce dTmax across THREADS */
17 shared double dTmax_local[THREADS];
18 /* ----- */
19
20
21 /* ----- */
22 /* Declare two new arrays of shared pointers, having same block sizes as
23 | * grid[][] and new_grid[][] - The dimension of each pointer can be only N+2
24 | * since each pointer only needs to point to the beginning of each row of
25 | * grid[][] or new_grid[][].
26 | * A single shared pointer can also be declared for temporary use in the
27 | * pointer swapping operation (we mean by shared arrays of pointers here
28 | * private arrays of pointers, each element pointing to a shared area)*/
29 shared [N+2] double *ptr[N+2], *new_ptr[N+2], *tmp_ptr;
30 /* ----- */
31
```

```

32 void initialize(void)
33 {
34     int j;
35
36     for( j=1; j<N+2; j++ )
37     {
38         grid[0][j] = 1.0;
39         new_grid[0][j] = 1.0;
40     }
41 }
42
43 int main(void)
44 {
45     struct timeval ts_st, ts_end;
46     double dTmax, dT, epsilon, max_time;
47     int finished, i, j, k, l;
48     double T;
49     int nr_iter;
50
51     if( MYTHREAD == 0 )
52         initialize();
53
54     /* ----- */
55     /* The next for loop is to initialize the pointers ptr[] and
56        ptr_new[] */
57     for( i=0; i<N+2; i++ )
58     {
59         ptr[i] = &grid[i][0];
60         new_ptr[i] = &new_grid[i][0];
61     }
62     /* ----- */
63
64     epsilon = 0.0001;
65     finished = 0;
66     nr_iter = 0;
67
68     upc_barrier;
69
70     gettimeofday( &ts_st, NULL );
71
72     do
73     {
74         dTmax = 0.0;
75         upc_forall( i=1; i<N+1; i++; i*THREADS/(N+2) )
76         {
77             for( j=1; j<N+1; j++ )
78             {
79                 /* ----- */
80                 /* this section has to be changed, to use ptr[][]
81                    and new_ptr[][] */
82                 T = 0.25 *
83                     ((ptr[i+1][j]) + (*ptr[i-1][j]) +
84                      (*ptr[i][j-1]) + (*ptr[i][j+1])); /* stencil */
85                 dT = T - ptr[i][j];
86                 new_ptr[i][j] = T;
87                 /* ----- */
88                 if( dTmax < fabs(dT) )
89                     dTmax = fabs(dT);
90             }
91         }

```

```

92
93     dTmax_local[MYTHREAD] = dTmax;
94     upc_barrier;
95     dTmax = dTmax_local[0];
96     for( i=1; i<THREADS; i++ )
97         if( dTmax < dTmax_local[i] )
98             dTmax = dTmax_local[i];
99     /* ----- */
100    /* here a barrier as to be introduced to avoid race
101       conditions over dTmax_local[i] */
102    upc_barrier;
103
104    if( dTmax < epsilon )
105        finished = 1;
106    else
107    {
108        /* ----- */
109        /* here swap the pointers ptr[] and new_ptr[] thru tmp_ptr
110       Note: a shared pivot pointer has to be declared */
111        for( k=0; k<N+2; k++ )
112        {
113            tmp_ptr = ptr[k];
114            ptr[k] = new_ptr[k];
115            new_ptr[k] = tmp_ptr;
116        }
117        /* ----- */
118
119        /* the barrier can be removed here since each thread has its
120       own private pointers to grid[][] and new_grid[][] whereas in
121       the previous case we wanted to wait till copying
122       new_grid[][] to grid[][] is completed by all threads */
123    }
124    nr_iter++;
125    } while( finished == 0 );
126
127    gettimeofday( &ts_end, NULL );
128
129    exectime[MYTHREAD] = ts_end.tv_sec + (ts_end.tv_usec / 1000000.0);
130    exectime[MYTHREAD] -= ts_st.tv_sec + (ts_st.tv_usec / 1000000.0);
131
132    upc_barrier;
133
134    if( MYTHREAD == 0 )
135    {
136        max_time = exectime[MYTHREAD];
137        for( i=1; i<THREADS; i++ )
138            if( max_time < exectime[i] )
139                max_time = exectime[i];
140        printf("%d iterations in %.3lf sec\n", nr_iter, max_time);
141    }
142
143    return 0;
144 }
145

```

## 4.6 - Performance boost using privatization

In the new version I improved the program performance by using privatization. Indeed, it has been shown that the speed of UPC memory accesses greatly differs on the type of access performed. The private memory access is the fastest. In the following program the different threads are working with new private pointers which are `*ptr_priv[]` and `*new_ptr_priv[]`. Their size is only the number of local rows which contain the data that corresponding to each threads, this is not the total number of rows.

To do this implementation I needed to adapt the previous program. I splitted the original `upc_forall(...)` loop into three blocks (l. 60-100) to separate the processing of the first row, middle part and last row. We need to split this loop because now threads are only working with private pointers which are not pointing to the global arrays. In other words, each thread can only accesses to a part of the global data.

This new optimization needed too the adaptation of other part of the program, as the pointer flipping (l. 117) for example.

```
1  #include <stdio.h>
2  #include <math.h>
3  #include <sys/time.h>
4  #include <upc_relaxed.h>
5
6  #define N 30
7
8
9  shared [N+2] double grid[N+2][N+2], new_grid[N+2][N+2];
10 shared double exectime[THREADS];
11 shared double dTmax_local[THREADS];
12
13 shared [N+2] double *ptr[N+2], *new_ptr[N+2], *tmp_ptr;
14 /* New private pointers */
15 double *ptr_priv[N+2], *new_ptr_priv[N+2], *tmp_ptr_priv;
16
17 void initialize(void)
18 {
19     int j;
20
21     for( j=1; j<N+2; j++ )
22     {
23         grid[0][j] = 1.0;
24         new_grid[0][j] = 1.0;
25     }
26 }
27
28 int main(void)
29 {
30     struct timeval ts_st, ts_end;
31     double dTmax, dT, epsilon, max_time;
32     int finished, i, j, k, l;
33     double T;
```



```
34     int nr_iter;
35
36     if( MYTHREAD == 0 )
37         initialize();
38
39
40     for( i=0; i<N+2; i++ )
41     {
42         ptr[i] = &grid[i][0];
43         new_ptr[i] = &new_grid[i][0];
44     }
45     /* to initialize the private pointers */
46     ptr_priv = &grid[MYTHREAD][0];
47     new_ptr_priv = &new_grid[MYTHREAD][0];
48
49     epsilon = 0.0001;
50     finished = 0;
51     nr_iter = 0;
52
53     upc_barrier;
54
55     gettimeofday( &ts_st, NULL );
```

```

56
57     do
58     {
59         dTmax = 0.0;
60         /* block 1 */
61         upc_forall( i=0; i<1; i++; i*THREADS/(N+2) )
62         {
63             for( j=0; j<1; j++ )
64             {
65                 T = 0.5 * ((*ptr_priv[i+1][j]) + (*ptr_priv[i][j+1])); /* stencil */
66                 dT = T - ptr_priv[i][j];
67                 new_ptr_priv[i][j] = T;
68
69                 if( dTmax < fabs(dT) )
70                     dTmax = fabs(dT);
71             }
72         }
73         /* block 2 */
74         upc_forall( i=1; i<N+1; i++; i*THREADS/(N+2) )
75         {
76             for( j=1; j<N+1; j++ )
77             {
78                 T = 0.25 *
79                     ((*ptr_priv[i+1][j]) + (*ptr_priv[i-1][j]) +
80                     (*ptr_priv[i][j-1]) + (*ptr_priv[i][j+1])); /* stencil */
81                 dT = T - ptr_priv[i][j];
82                 new_ptr_priv[i][j] = T;
83
84                 if( dTmax < fabs(dT) )
85                     dTmax = fabs(dT);
86             }
87         }
88
89         /* block 3 */
90         upc_forall( i=N; i<N+2; i++; i*THREADS/(N+2) )
91         {
92             for( j=N; j<N+2; j++ )
93             {
94                 T = 0.5 * ((*ptr_priv[i-1][j]) + (*ptr_priv[i][j-1])); /* stencil */
95                 dT = T - ptr_priv[i][j];
96                 new_ptr_priv[i][j] = T;
97
98                 if( dTmax < fabs(dT) )
99                     dTmax = fabs(dT);
100             }
101         }
102
103         dTmax_local[MYTHREAD] = dTmax;
104         upc_barrier;
105         dTmax = dTmax_local[0];
106         for( i=1; i<THREADS; i++ )
107             if( dTmax < dTmax_local[i] )
108                 dTmax = dTmax_local[i];
109         upc_barrier;

```

```
110
111     if( dTmax < epsilon )
112     |     finished = 1;
113     else
114     {
115         for( k=0; k<N+2; k++ )
116         {
117             /* Pointer flipping for private pointers */
118             tmp_ptr_priv    = ptr_priv[k];
119             ptr_priv[k]     = new_ptr_priv[k];
120             new_ptr_priv[k] = tmp_ptr_priv;
121             /* --- */
122             tmp_ptr    = ptr[k];
123             ptr[k]     = new_ptr[k];
124             new_ptr[k] = tmp_ptr;
125         }
126     }
127     nr_iter++;
128 } while( finished == 0 );
129
130 gettimeofday( &ts_end, NULL );
131
132 exectime[MYTHREAD] = ts_end.tv_sec + (ts_end.tv_usec / 1000000.0);
133 exectime[MYTHREAD] -= ts_st.tv_sec + (ts_st.tv_usec / 1000000.0);
134
135 upc_barrier;
136
137 if( MYTHREAD == 0 )
138 {
139     max_time = exectime[MYTHREAD];
140
141     for( i=1; i<THREADS; i++ )
142     |     if( max_time < exectime[i] )
143     |     |     max_time = exectime[i];
144     printf("%d iterations in %.3lf sec\n", nr_iter, max_time);
145 }
146
147 return 0;
148 }
149
```

## 4.7 - Dynamic problem size

In this last version, I implemented the dynamic allocating shared memory during execution to make the problem size dynamic. In other word, user can specifying N at runtime. As we can see in the program, now all pointer are declare without any dependencies to N.

```
1  #include <stdio.h>
2  #include <math.h>
3  #include <sys/time.h>
4  #include <upc_relaxed.h>
5
6  // #define N 30
7  int N;
8
9  shared [] double grid[][], new_grid[][];
10 shared double exectime[THREADS];
11 shared double dTmax_local[THREADS];
12
13 shared [] double *ptr[], *new_ptr[], *tmp_ptr;
14 /* New private pointers */
15 double *ptr_priv[], *new_ptr_priv[], *tmp_ptr_priv;
16
17 void initialize(void)
18 {
19     int j;
20
21     for( j=1; j<N+2; j++ )
22     {
23         grid[0][j] = 1.0;
24         new_grid[0][j] = 1.0;
25     }
26 }
27
28 int main(int argc, char** argv)
29 {
30     struct timeval ts_st, ts_end;
31     double dTmax, dT, epsilon, max_time;
32     int finished, i, j, k, l;
33     double T;
```

```
34     int nr_iter;
35
36     if(argc != 2) {
37         if(MYTHREAD == 0)
38             printf("Erreur : Un parametre est attendu. %s\n", argv[0]);
39         return -1;
40     }
41
42     /* to get N */
43     if(argv[1] <= 0) {
44         if(MYTHREAD == 0)
45             printf("Erreur : N should be bigger than 0 \n");
46         return -2;
47     }
48     N = argv[1];
49
50     /* dynamic allocation */
51     if( MYTHREAD == 0 ) {
52         // grid & new_grid
53         grid = (shared[N+2] double *) upc_alloc((N+2)*(N+2) * sizeof(double));
54         new_grid = (shared[N+2] double *) upc_alloc((N+2)*(N+2) * sizeof(double));
55         // ptr & new_ptr
56         ptr = (shared[] double *) upc_alloc((N+2) * sizeof(double));
57         new_ptr = (shared[] double *) upc_alloc((N+2) * sizeof(double));
58     }
59     upc_barrier;
60     // ptr_priv & new_ptr_priv
61     ptr_priv = (shared[] double *) upc_alloc((N+2) * sizeof(double));
62     new_ptr_priv = (shared[] double *) upc_alloc((N+2) * sizeof(double));
63     upc_barrier;
64
```



```
65
66     if( MYTHREAD == 0 )
67         initialize();
68
69     for( i=0; i<N+2; i++ )
70     {
71         ptr[i] = &grid[i][0];
72         new_ptr[i] = &new_grid[i][0];
73     }
74     /* to initialize the private pointers */
75     ptr_priv = &grid[MYTHREAD][0];
76     new_ptr_priv = &new_grid[MYTHREAD][0];
77
78     epsilon = 0.0001;
79     finished = 0;
80     nr_iter = 0;
81
82     upc_barrier;
83
84     gettimeofday( &ts_st, NULL );
85
86     do
87     {
88         dTmax = 0.0;
89         /* block 1 */
90         upc_forall( i=0; i<1; i++; i*THREADS/(N+2) )
91         {
92             for( j=0; j<1; j++ )
93             {
94                 T = 0.5 * ((*ptr_priv[i+1][j]) + (*ptr_priv[i][j+1])); /* stencil */
95                 dT = T - ptr_priv[i][j];
96                 new_ptr_priv[i][j] = T;
```

```

97
98     if( dTmax < fabs(dT) )
99         dTmax = fabs(dT);
100
101 }
102
103 /* block 2 */
104 upc_forall( i=1; i<N+1; i++; i*THREADS/(N+2) )
105 {
106     for( j=1; j<N+1; j++ )
107     {
108         T = 0.25 *
109             (( *ptr_priv[i+1][j] ) + ( *ptr_priv[i-1][j] ) +
110              ( *ptr_priv[i][j-1] ) + ( *ptr_priv[i][j+1] )); /* stencil */
111         dT = T - ptr_priv[i][j];
112         new_ptr_priv[i][j] = T;
113
114         if( dTmax < fabs(dT) )
115             dTmax = fabs(dT);
116     }
117
118 /* block 3 */
119 upc_forall( i=N; i<N+2; i++; i*THREADS/(N+2) )
120 {
121     for( j=N; j<N+2; j++ )
122     {
123         T = 0.5 * (( *ptr_priv[i-1][j] ) + ( *ptr_priv[i][j-1] )); /* stencil */
124         dT = T - ptr_priv[i][j];
125         new_ptr_priv[i][j] = T;
126
127         if( dTmax < fabs(dT) )
128             dTmax = fabs(dT);
129     }
130 }
```

```
130
131     dTmax_local[MYTHREAD] = dTmax;
132     upc_barrier;
133     dTmax = dTmax_local[0];
134     for( i=1; i<THREADS; i++ )
135         if( dTmax < dTmax_local[i] )
136             dTmax = dTmax_local[i];
137
138     upc_barrier;
139
140     if( dTmax < epsilon )
141         finished = 1;
142     else
143     {
144         for( k=0; k<N+2; k++ )
145         {
146             /* Pointer flipping for private pointers */
147             tmp_ptr_priv = ptr_priv[k];
148             ptr_priv[k] = new_ptr_priv[k];
149             new_ptr_priv[k] = tmp_ptr_priv;
150             /* --- */
151             tmp_ptr = ptr[k];
152             ptr[k] = new_ptr[k];
153             new_ptr[k] = tmp_ptr;
154         }
155     }
156     nr_iter++;
157 } while( finished == 0 );
158
159 gettimeofday( &ts_end, NULL );
```

```
160
161     exectime[MYTHREAD] = ts_end.tv_sec + (ts_end.tv_usec / 1000000.0);
162     exectime[MYTHREAD] -= ts_st.tv_sec + (ts_st.tv_usec / 1000000.0);
163
164     upc_barrier;
165
166     if( MYTHREAD == 0 )
167     {
168         max_time = exectime[MYTHREAD];
169         for( i=1; i<THREADS; i++ )
170             if( max_time < exectime[i] )
171                 max_time = exectime[i];
172         printf("%d iterations in %.3lf sec\n", nr_iter, max_time);
173     }
174
175     return 0;
176 }
177
178
```